

HW1: Tung Pham

Collaboration Statement: This assignment was completed with collaboration on concept discussion, code-debugging and idea exchange from Arman Muratbayev, Avtar Rekhi and Joseph Hadidjojo

Total hours of human time: approx. 10-15 hours

Total hours of machine time: approx. 10-15 hrs

Links: [HW1 instructions] [Course collaboration policy]

Contents

1a: Figure	2
1b: Figure	3
2a: Solution	4
2b: Solution	4
3a: Concept Question	5

1a: Figure

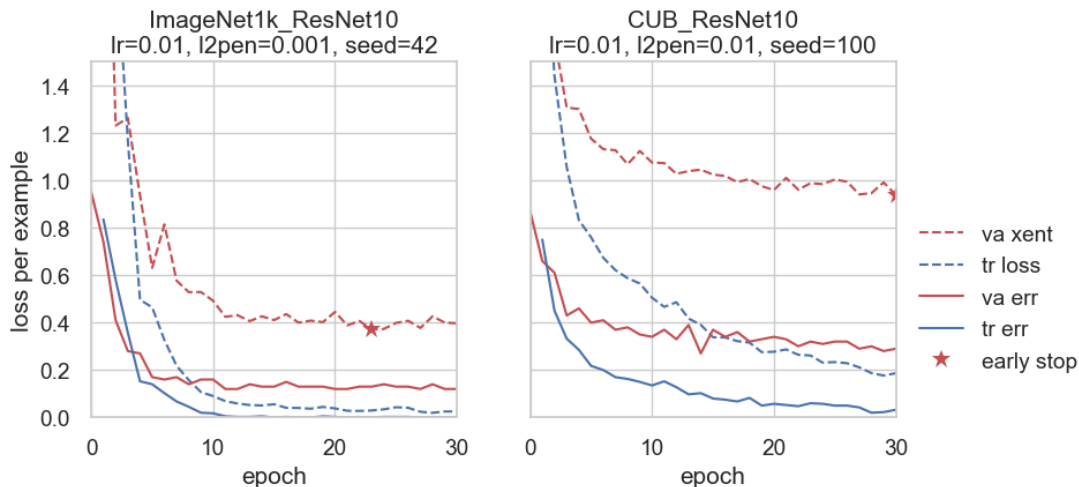


Fig. 1a: The two graphs display the curve trends of the losses and error of two different models that was trained on 2 different dataset. They all converge nicely and there is no sign of overfitting as there isn't any upward trends in the validation loss or error while there is a downwards trends or decrease in the training loss or error which indicates that both models are fitting the training set very well. The model that was trained on ImageNet1k decrease rather quicker than CUB and flattens out and not improving until early-stopping occurs. However, we can see that the validation cross-entropy value is converging at a rather higher value compares to error which indicates that the model is not very confident in its selection. As for model that was trained on CUB dataset, the training loss is rather slower to converge as we can still see sign of decreasing after we've finished the whole epochs. However, for the losses and validation cross-entropy value is comparatively fast to converge to the model trained ImageNet1k dataset. There is also a sign of unconfidence detected in this model as well since we're seeing a significant larger value of cross-entropy compares to the training loss at each epoch for this model. In conclusion, it seems that the model trained with ImageNet1k generalize better than the model that trained with CUB as the validation loss and training loss are closer to each other. During experimenting, early stopping helps boost the training process significantly as it reduce the need for further computation on some instance which reduce the time to conduct grid search on the hyperparameters. I've test with multiple a stronger L2 penalty but the effect on the loss value seems unnoticeable.

1b: Figure

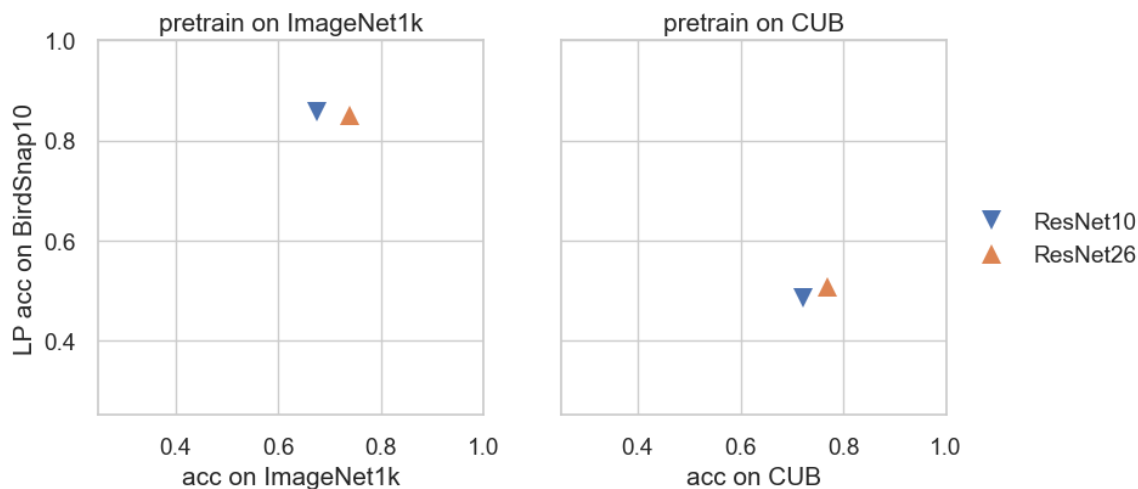


Fig. 1b: From the plot, we can see that the model pretrained on the ImageNet1k perform better than pretrained on CUB. We can also see that the result obtained from our model exceed the top-1 accuracy on Pytorchcv model using ResNet10 on ImageNet1k dataset. This confirms the observation of Fang et. al. in his "Does progress on ImageNet transfer to real-world datasets" as we can see that this models perform exceedingly well on the ImageNet dataset, however, when transferring to another dataset like CUB, the performance of the model significantly drop. Another note is that on both of our graph, using a larger model indeed improve the accuracy which confirms that a larger model with more layers and capacity indeed improve our performance which support Fang's claim that the model complexity can help in transfer tasks in some cases.

2a: Solution

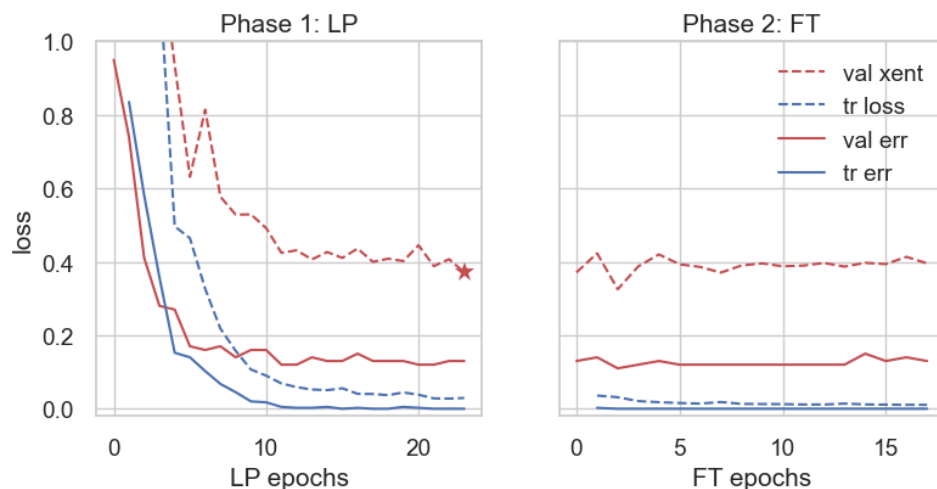


Fig. 2a: In this experiment, a grid search was conducted to find the best hyperparameters for the fine-tuning phase. However, because we're keeping the hyperparameters from previous experiments which fits the data too well, we can see little to no gain on the second phase of the training. In the first phase, we can see the clear sign of improvements gained from the Linear Probing where the loss and error rates drop quickly into convergence around epoch 10. However, in the second phase, as we've hope that the Fine tuning on the last 3 layers would help further reduce the loss and error rate down, we didn't see this result from the graph. In the second phase, the losses and error remains constants throughout the training with minimal to no drop on any benchmarks. It almost have an upward trend which is an indication of overfitting of the second model. It could potentially because in the first phase, the model has fit too well on the data cause limited to no room in improving the scores on the second phase and potentially causing overfitting. This indicates a failure in our implementation to achieve the improvement of this approach.

2b: Solution

method	test acc
-----	-----
LP	0.857
LP-then-FT	0.850

Here we can see that LP-then-FT is actually not improving our test accuracy but instead made it worsen. This is due to the fact that it was not improving the loss or errors rate during training and potentially cause overfitting as we kept on training. This could be because during Linear Probing, the model has already fit too well on the dataset which cause limited space for improvements during Fine Tuning and potentially cause overfitting on this phase. As seen in the graph before, the graph have an almost upward trend which potentially is the sign of overfitting and cause of this degradation in accuracy on unseen data.

3a: Concept Question

Given that z_i represent the logits vector for example i , we can say the following:

$$L_{z_i=x_i+b}$$

where b is the vector of bias parameters of last layer and x_i is the input vector which is assumed to have the size of $64 \times 3 \times 244 \times 244$. This means that x_i is an input of 64 images and each image has 3 channel. The size of each image is 244×244 which makes i is in the range of 1 to 64.

For simplicity, let's assume that there is no activate function being used here for forward propagation. Since in our model `predict_proba` we call `softmax` on the forward, we can come up with the following formula for softmax:

$$-\log \left(\frac{\exp(z_{i,y_i})}{\sum_{c=1}^C \exp(z_{i,c})} \right)$$

Our L2 penalty can be written as the sum of square of all the weights which is defined as follow:

$$\lambda \sum_{j=1}^d \sum_{c=1}^C w_{c,j}^2$$

Where λ is the L2 penalty magnitude, C is the number of classes for classification which is 10 and weight w which has the size d by C where d is the dimension of the previous layer `AveragePool2D` of 512 and C is 10.

To finalize our loss function, we just need to take the L2 penalty and add to the normalized softmax across the batch which give us the following:

$$L_{\text{total}} = \frac{1}{B} \sum_{i=1}^B -\log \left(\frac{\exp(z_{i,y_i})}{\sum_{c=1}^C \exp(z_{i,c})} \right) + \lambda \sum_{j=1}^d \sum_{c=1}^C w_{c,j}^2$$